

“Box pruning revisited” – an optimization project by Pierre Terdiman – 2017

Part 11 – the last branch

The SIMD version gave us a good speedup. The best scalar version was also quite successful, thanks to the removal of several hard-to-predict branches.

Looking at the remaining code with this in mind, one more branch jumps to the eyes, just before the SIMD overlap test:

```
if(Index0!=Index1)
{
```

This ensures that we never test a box against itself, but I don't quite remember why it's there. The way the code is written, it cannot happen. So this test can just be removed. Since it never happens, the branch is likely to be always correctly predicted, and thus the expected gains are small, if any.

But still, worth doing - it's just one line.

And the results are actually interesting:

Home PC:

Unsafe version:

Complete test (brute force): found 11811 intersections in 781447 K-cycles.

15224 K-cycles.

14386 K-cycles.

14633 K-cycles.

14387 K-cycles.

14376 K-cycles.

14392 K-cycles.

14376 K-cycles.

14373 K-cycles.

14579 K-cycles.

14400 K-cycles.

14374 K-cycles.

14388 K-cycles.

14373 K-cycles.

14372 K-cycles.

14471 K-cycles.

14404 K-cycles.

Complete test (box pruning): found 11715 intersections in 14372 K-cycles.

Safe version:

Complete test (brute force): found 11811 intersections in 781701 K-cycles.

15306 K-cycles.

14533 K-cycles.

14513 K-cycles.

14731 K-cycles.

14552 K-cycles.

14512 K-cycles.

14528 K-cycles.

14514 K-cycles.

14514 K-cycles.

14528 K-cycles.

14535 K-cycles.

14513 K-cycles.

14526 K-cycles.

14515 K-cycles.

14515 K-cycles.

14621 K-cycles.

Complete test (box pruning): found 11811 intersections in 14512 K-cycles.

Office PC:

Unsafe version:

Complete test (brute force): found 11811 intersections in 824867 K-cycles.

13618 K-cycles.

12900 K-cycles.

12573 K-cycles.

12957 K-cycles.

12570 K-cycles.

12911 K-cycles.

12572 K-cycles.

12573 K-cycles.

12588 K-cycles.

13153 K-cycles.

13447 K-cycles.

13212 K-cycles.

13429 K-cycles.

13214 K-cycles.

13527 K-cycles.

13229 K-cycles.

Complete test (box pruning): found 11715 intersections in 12570 K-cycles.

Safe version:

Complete test (brute force): found 11811 intersections in 816240 K-cycles.

13227 K-cycles.

13013 K-cycles.

12621 K-cycles.

14329 K-cycles.

12624 K-cycles.

13009 K-cycles.

13533 K-cycles.

13171 K-cycles.

12881 K-cycles.

13181 K-cycles.

13265 K-cycles.

13248 K-cycles.

13245 K-cycles.

13242 K-cycles.

13267 K-cycles.

12611 K-cycles.

Complete test (box pruning): found 11811 intersections in 12611 K-cycles.

The gains are summarized here:

Home PC	Timings (K-Cycles)	Delta (K-Cycles)	Speedup	Overall X factor
(Version1)	(101662)			
Version2 - base	98822	0	0%	1.0
Version3	93138	~5600	~5%	~1.06
Version4	81834	~11000	~12%	~1.20
Version5	78140	~3600	~4%	~1.26
Version6a	60579	~17000	~22%	~1.63
Version6b	41605	~18000	~31%	~2.37
(Version7)	(40906)	-	-	-
(Version8)	(31383)	(~10000)	(~24%)	(~3.14)
Version9a	34486	~7100	~17%	~2.86
Version9b - unsafe	32477	~2000	~5%	~3.04
Version9b - safe	32565	~1900	~5%	~3.03
Version9c - unsafe	16223	~16000	~50%	~6.09
Version9c - safe	14802	~17000	~54%	~6.67
(Version10)	(16667)	-	-	-
Version11 - unsafe	14372	~1800	~11%	~6.87
Version11 - safe	14512	~200	~2%	~6.80

Office PC	Timings (K-Cycles)	Delta (K-Cycles)	Speedup	Overall X factor
(Version1)	(96203)			
Version2 - base	92885	0	0%	1.0
Version3	88352	~4500	~5%	~1.05
Version4	77156	~11000	~12%	~1.20
Version5	73778	~3300	~4%	~1.25
Version6a	58451	~15000	~20%	~1.58
Version6b	45634	~12000	~21%	~2.03
(Version7)	(43987)	-	-	-
(Version8)	(29083)	(~16000)	(~36%)	(~3.19)
Version9a	31864	~13000	~30%	~2.91
Version9b - unsafe	15097	~16000	~52%	~6.15
Version9b - safe	15116	~16000	~52%	~6.14
Version9c - unsafe	12707	~2300	~15%	~7.30
Version9c - safe	12562	~2500	~16%	~7.39
(Version10)	(15648)	-	-	-
Version11 - unsafe	12570	~100	~1%	~7.38
Version11 - safe	12611	-	-	~7.36

Version 11 is compared to version 9c here.

No difference on the office PC, but on the home PC the safe & unsafe versions are suddenly the same speed. Hmmm. We'll never solve that particular mystery then. Oh well. All is good now at least.

Now what? Are we done?

We are about 7x faster than when we started, which is pretty good!

...but not quite the order of magnitude we wanted to reach.

At this point there isn't much left in the inner loop, at least on the C++ side of things. So let's check the new disassembly. The inner loop we're interested in is:

```
while(BoxListX[Index1].mMinX<=MaxLimit)
{
    SIMD_OVERLAP_TEST(BoxListYZ[Index1])
    pairs.Add(RIndex0).Add(Remap[Index1]);

    Index1++;
}
```

And the corresponding disassembly is simple enough to isolate:

```
//SIMD_OVERLAP_TEST(BoxListYZ[Index1])
003E2FB3 cmpltps   xmm1,xmmword ptr [edi]
003E2FB7 movmskps  eax,xmm1
003E2FBA cmp      eax,0Ch
003E2FBD jne      CompleteBoxPruning+2D1h (03E3011h)

//pairs.Add(RIndex0).Add(Remap[Index1]);
003E2FBF mov      eax,dword ptr [esi+4]
003E2FC2 cmp      eax,dword ptr [esi]
003E2FC4 jne      CompleteBoxPruning+28Fh (03E2FCFh)
003E2FC6 push     1
003E2FC8 mov      ecx,esi
003E2FCA call     IceCore::Container::Resize (03E1350h)
003E2FCF mov      ecx,dword ptr [esi+4]
003E2FD2 mov      eax,dword ptr [esi+8]
003E2FD5 mov      edx,dword ptr [esp+34h]
003E2FD9 mov      dword ptr [eax+ecx*4],edx
003E2FDC inc      dword ptr [esi+4]
003E2FDF mov      edx,dword ptr [esp+20h]
003E2FE3 mov      eax,dword ptr [esi+4]
003E2FE6 mov      ecx,dword ptr [edx]
003E2FE8 mov      dword ptr [esp+30h],ecx
003E2FEC cmp      eax,dword ptr [esi]
003E2FEE jne      CompleteBoxPruning+2B9h (03E2FF9h)
003E2FF0 push     1
003E2FF2 mov      ecx,esi
003E2FF4 call     IceCore::Container::Resize (03E1350h)
003E2FF9 mov      ecx,dword ptr [esi+4]
003E2FFC mov      eax,dword ptr [esi+8]
```

```

003E2FFF mov     edx,dword ptr [esp+30h]
003E3003 mov     dword ptr [eax+ecx*4],edx
003E3006 inc     dword ptr [esi+4]
003E3009 mov     ecx,dword ptr [esp+28h]
003E300D mov     edx,dword ptr [esp+20h]

// while(BoxListX[Index1].mMinX<=MaxLimit)
003E3011 movss   xmm0,dword ptr [esp+38h]
003E3017 movaps  xmm1,xmmword ptr [esp+40h]
003E301C add     ecx,8
003E301F add     edx,4
003E3022 add     edi,10h
003E3025 comiss  xmm0,dword ptr [ecx]
003E3028 mov     dword ptr [esp+20h],edx
003E302C mov     dword ptr [esp+28h],ecx
003E3030 jae     CompleteBoxPruning+273h (03E2FB3h)

```

I added two blank lines and color-coded the whole thing to clearly delimitate three blocks of code.

Roughly, the **first block** is the SIMD overlap test, the **second block** is for writing the pair indices when an overlap is found, and the **last block** is the while loop (easily identified by the *comiss* instruction).

It's... not great.

Where do I start?

In the first two lines, since the *xmmword ptr [edi]* is the actual load of *BoxListYZ[Index1]*, writing it this way means that *xmm1* contains our constant box for the loop (the preloaded *Box0YZ* in the C++ code):

```

003E2FB3 cmpltps  xmm1,xmmword ptr [edi]
003E2FB7 movmskps eax,xmm1

```

But the *cmpltps* instruction is always going to destroy the contents of *xmm1*. So this register will need to be reloaded for the next iteration. And indeed, that's what happens in the third block:

```

003E3017 movaps  xmm1,xmmword ptr [esp+40h]

```

We also confirm that *edi* is the *BoxListYZ* array, thanks to:

```

003E3022 add     edi,10h

```

10h == 16 == sizeof(SIMD_AABB_YZ), makes sense.

Question 1: why write the code in such a way that it destroys the constant box?

Question 2: why reload the constant box from the stack? Why not keeping it in a register? The whole loop only uses two *xmm* registers (!) so it's not like there's a shortage of them.

In the third block, we have this:

```
003E301C add    ecx,8
```

And then:

```
003E3025 comiss xmm0,dword ptr [ecx]
```

Which means that *ecx* is the *BoxListX* array (with `sizeof(SIMD_AABB_X)==8`), and *dword ptr [ecx]* (a scalar load) is the assembly for "*BoxListX[Index1].mMinX*" in the C++ code.

Makes sense.

But then it also means that *xmm0* is "*MaxLimit*", which gets reloaded just before from the stack:

```
003E3011 movss   xmm0,dword ptr [esp+38h]
```

Doesn't make sense.

Question 3: Why do you reload it from the stack each time? It's a constant for the whole loop. There are free available *xmm* registers. Just put it in *xmm2*, keep it there, and the "*movss*" instruction above just vanishes.

In the third block again, we have this:

```
003E301F add    edx,4
...
003E3028 mov     dword ptr [esp+20h],edx
003E302C mov     dword ptr [esp+28h],ecx
```

We saw *ecx* before, it's the *BoxListX* array.

Question 4: why do you push it to the stack? Why don't you keep it in a register?

Then there's *edx*, incremented by 4 and also pushed to the stack. If you look for *[esp+20h]* to check where we read it again, you find two places:

```
003E2FDF mov    edx,dword ptr [esp+20h]
```

And:

```
003E300D mov    edx,dword ptr [esp+20h]
```

The first place happens after the first "push back" call, and corresponds to this C++ code:

```
Add(Remap[Index1])
```

The second place happens after the second "push back" call, and simply restores the *edx* register, which is not actually used when the SIMD-overlap test doesn't find an overlap.

You guessed it, *edx* is *Index1*, or rather *Remap[Index1]* as you can see from these lines in the second block:

```
003E2FDF mov    edx,dword ptr [esp+20h]
...
003E2FE6 mov    ecx,dword ptr [edx]
```

Right. Ok. But then...

Question 5: why do you save *edx* to the stack all the time (address 003E3028)? Why don't you just save and restore it only when an overlap actually occurs?

Question 6: same question for *ecx*.

As for the second block, it looks fine overall. Still, I know some people will ask this one:

Question 7: why do you use a custom vector class instead of *std::vector*?

Ok, let's quickly deal with that one first, it won't take long. Replace the custom container with *std::vector* and you get:

Home PC:

Unsafe:

Complete test (brute force): found 11811 intersections in 851569 K-cycles.

19727 K-cycles.

18913 K-cycles.

18881 K-cycles.

18879 K-cycles.

18908 K-cycles.

18901 K-cycles.

18891 K-cycles.

18882 K-cycles.

18878 K-cycles.

19110 K-cycles.

18902 K-cycles.

18901 K-cycles.

18881 K-cycles.

18895 K-cycles.

18904 K-cycles.

18882 K-cycles.

Complete test (box pruning): found 11715 intersections in 18878 K-cycles.

Office PC:

Unsafe:

Complete test (brute force): found 11811 intersections in 888075 K-cycles.

19200 K-cycles.

18307 K-cycles.

18584 K-cycles.

18729 K-cycles.

18306 K-cycles.

18647 K-cycles.

18571 K-cycles.

18306 K-cycles.

18767 K-cycles.

18551 K-cycles.

18420 K-cycles.

18659 K-cycles.

18530 K-cycles.

18365 K-cycles.

18491 K-cycles.

18310 K-cycles.

Complete test (box pruning): found 11715 intersections in 18306 K-cycles.

That is:

Home PC	Timings
Version11 - unsafe - ICE container	14372
Version11 - unsafe - std::vector	18878

Office PC	Timings
Version11 - unsafe - ICE container	12570
Version11 - unsafe - std::vector	18306

Just look at the numbers. That's why I'm not using *std::vector*. It has never been able to implement a simple "push_back" correctly. It's 2017 and *std::vector* still cannot match the first basic C++ class I wrote back in 1999.

There is no need to tell me about *STLPort* or *EASTL* or whatever: I've heard it all. If I cannot rely on the version included with *Visual Studio*, if I have to switch to some external library, I'd rather use my own. At least I know it's properly implemented.

Just for "fun", with the custom container the [second block](#) above is about 27 lines of code. Right?

Ok, now take a deep breath, here's the equivalent second block with *std::vector*:

```
011836B0 mov     ecx,dword ptr [esi+4]
011836B3 lea     eax,[esp+28h]
011836B7 cmp     eax,ecx
011836B9 jae     CompleteBoxPruning+316h (01183736h)
011836BB mov     edi,dword ptr [esi]
011836BD cmp     edi,eax
011836BF ja     CompleteBoxPruning+316h (01183736h)
011836C1 mov     edx,dword ptr [esi+8]
011836C4 sub     eax,edi
011836C6 sar     eax,2
011836C9 mov     dword ptr [esp+40h],eax
011836CD cmp     ecx,edx
011836CF jne     CompleteBoxPruning+302h (01183722h)
011836D1 mov     eax,edx
011836D3 sub     eax,ecx
011836D5 sar     eax,2
011836D8 cmp     eax,1
011836DB jae     CompleteBoxPruning+302h (01183722h)
```

```

011836DD sub    ecx,edi
011836DF sar    ecx,2
011836E2 mov    eax,3FFFFFFFh
011836E7 sub    eax,ecx
011836E9 cmp    eax,1
011836EC jb     CompleteBoxPruning+4F9h (01183919h)
011836F2 inc    ecx
011836F3 sub    edx,edi
011836F5 sar    edx,2
011836F8 mov    dword ptr [esp+3Ch],ecx
011836FC mov    ecx,edx
011836FE shr    ecx,1
01183700 mov    eax,3FFFFFFFh
01183705 sub    eax,ecx
01183707 cmp    eax,edx
01183709 jae    CompleteBoxPruning+2EFh (0118370Fh)
0118370B xor    edx,edx
0118370D jmp    CompleteBoxPruning+2F1h (01183711h)
0118370F add    edx,ecx
01183711 cmp    edx,dword ptr [esp+3Ch]
01183715 mov    ecx,esi
01183717 cmovb  edx,dword ptr [esp+3Ch]
0118371C push   edx
0118371D call   std::vector<unsigned int,std::allocator<unsigned int> >::_Reallocate (011828F0h)
01183722 mov    edx,dword ptr [esi+4]
01183725 test   edx,edx
01183727 je     CompleteBoxPruning+37Dh (0118379Dh)
01183729 mov    eax,dword ptr [esi]
0118372B mov    ecx,dword ptr [esp+40h]
0118372F mov    eax,dword ptr [eax+ecx*4]
01183732 mov    dword ptr [edx],eax
01183734 jmp    CompleteBoxPruning+37Dh (0118379Dh)
01183736 mov    edx,dword ptr [esi+8]
01183739 cmp    ecx,edx
0118373B jne    CompleteBoxPruning+370h (01183790h)
0118373D mov    eax,edx
0118373F sub    eax,ecx
01183741 sar    eax,2
01183744 cmp    eax,1
01183747 jae    CompleteBoxPruning+370h (01183790h)
01183749 mov    edi,dword ptr [esi]
0118374B sub    ecx,edi
0118374D sar    ecx,2
01183750 mov    eax,3FFFFFFFh
01183755 sub    eax,ecx
01183757 cmp    eax,1
0118375A jb     CompleteBoxPruning+4F9h (01183919h)
01183760 inc    ecx

```

```

01183761 sub    edx,edi
01183763 sar    edx,2
01183766 mov    dword ptr [esp+40h],ecx
0118376A mov    ecx,edx
0118376C shr    ecx,1
0118376E mov    eax,3FFFFFFFh
01183773 sub    eax,ecx
01183775 cmp    eax,edx
01183777 jae    CompleteBoxPruning+35Dh (0118377Dh)
01183779 xor    edx,edx
0118377B jmp    CompleteBoxPruning+35Fh (0118377Fh)
0118377D add    edx,ecx
0118377F cmp    edx,dword ptr [esp+40h]
01183783 mov    ecx,esi
01183785 cmovb  edx,dword ptr [esp+40h]
0118378A push   edx
0118378B call  std::vector<unsigned int,std::allocator<unsigned int> >::_Reallocate (011828F0h)
01183790 mov    eax,dword ptr [esi+4]
01183793 test   eax,eax
01183795 je     CompleteBoxPruning+37Dh (0118379Dh)
01183797 mov    ecx,dword ptr [esp+44h]
0118379B mov    dword ptr [eax],ecx
0118379D add    dword ptr [esi+4],4
011837A1 mov    ecx,dword ptr [esi+4]
011837A4 mov    eax,dword ptr [esp+14h]
011837A8 cmp    eax,ecx
011837AA jae    CompleteBoxPruning+405h (01183825h)
011837AC mov    edi,dword ptr [esi]
011837AE cmp    edi,eax
011837B0 ja     CompleteBoxPruning+405h (01183825h)
011837B2 mov    edx,dword ptr [esi+8]
011837B5 sub    eax,edi
011837B7 sar    eax,2
011837BA mov    dword ptr [esp+3Ch],eax
011837BE cmp    ecx,edx
011837C0 jne    CompleteBoxPruning+3F3h (01183813h)
011837C2 mov    eax,edx
011837C4 sub    eax,ecx
011837C6 sar    eax,2
011837C9 cmp    eax,1
011837CC jae    CompleteBoxPruning+3F3h (01183813h)
011837CE sub    ecx,edi
011837D0 sar    ecx,2
011837D3 mov    eax,3FFFFFFFh
011837D8 sub    eax,ecx
011837DA cmp    eax,1
011837DD jb     CompleteBoxPruning+4F9h (01183919h)
011837E3 inc    ecx

```

```

011837E4 sub     edx,edi
011837E6 sar     edx,2
011837E9 mov     dword ptr [esp+40h],ecx
011837ED mov     ecx,edx
011837EF shr     ecx,1
011837F1 mov     eax,3FFFFFFFh
011837F6 sub     eax,ecx
011837F8 cmp     eax,edx
011837FA jae     CompleteBoxPruning+3E0h (01183800h)
011837FC xor     edx,edx
011837FE jmp     CompleteBoxPruning+3E2h (01183802h)
01183800 add     edx,ecx
01183802 cmp     edx,dword ptr [esp+40h]
01183806 mov     ecx,esi
01183808 cmovb   edx,dword ptr [esp+40h]
0118380D push    edx
0118380E call    std::vector<unsigned int,std::allocator<unsigned int> >::_Reallocate (011828F0h)
01183813 mov     ecx,dword ptr [esi+4]
01183816 test    ecx,ecx
01183818 je      CompleteBoxPruning+46Eh (0118388Eh)
0118381A mov     eax,dword ptr [esi]
0118381C mov     edx,dword ptr [esp+3Ch]
01183820 mov     eax,dword ptr [eax+edx*4]
01183823 jmp     CompleteBoxPruning+46Ch (0118388Ch)
01183825 mov     edx,dword ptr [esi+8]
01183828 cmp     ecx,edx
0118382A jne     CompleteBoxPruning+463h (01183883h)
0118382C mov     eax,edx
0118382E sub     eax,ecx
01183830 sar     eax,2
01183833 cmp     eax,1
01183836 jae     CompleteBoxPruning+45Fh (0118387Fh)
01183838 mov     edi,dword ptr [esi]
0118383A sub     ecx,edi
0118383C sar     ecx,2
0118383F mov     eax,3FFFFFFFh
01183844 sub     eax,ecx
01183846 cmp     eax,1
01183849 jb      CompleteBoxPruning+4F9h (01183919h)
0118384F inc     ecx
01183850 sub     edx,edi
01183852 sar     edx,2
01183855 mov     dword ptr [esp+40h],ecx
01183859 mov     ecx,edx
0118385B shr     ecx,1
0118385D mov     eax,3FFFFFFFh
01183862 sub     eax,ecx
01183864 cmp     eax,edx

```

```

01183866 jae      CompleteBoxPruning+44Ch (0118386Ch)
01183868 xor      edx,edx
0118386A jmp      CompleteBoxPruning+44Eh (0118386Eh)
0118386C add      edx,ecx
0118386E cmp      edx,dword ptr [esp+40h]
01183872 mov      ecx,esi
01183874 cmovb    edx,dword ptr [esp+40h]
01183879 push     edx
0118387A call     std::vector<unsigned int,std::allocator<unsigned int> >::_Reallocate (011828F0h)
0118387F mov      eax,dword ptr [esp+14h]
01183883 mov      ecx,dword ptr [esi+4]
01183886 test     ecx,ecx
01183888 je       CompleteBoxPruning+46Eh (0118388Eh)
0118388A mov      eax,dword ptr [eax]
0118388C mov      dword ptr [ecx],eax
0118388E mov      edx,dword ptr [esp+2Ch]
01183892 mov      ecx,dword ptr [esp+14h]
01183896 add      dword ptr [esi+4],4

```

That's **180** lines of code. Instead of **27**. It's just ridiculous.

That's all I have to say about *std::vector*.

Now that this interlude is over, let's go back to the sane version presented before.

The most commonly executed codepath (when an overlap is not found) has 13 instructions (it corresponds to the **first block** + the **last block**). And for 13 instructions we raised 6 potential issues. That's a lot of questions for such a small piece of code.

Of course, we might be missing something. The compiler might know more than us. Our analysis might be too naive.

Perhaps.

Perhaps not.

How do you answer that?

As before: the scientific way. We observed. Now we need to come up with theories and tests to discern between facts and fiction.

Theory: the compiler is not recent enough. A more recent compiler would produce better code.

Yeah. I heard that one before.... for every version of *Visual Studio* since VC6... But fair enough, it's easy to test. In fact let's try both VC10 (the SSE code generation has changed quite a bit between VC10 and VC11), and then VC14. VC10 gives:

```
00BE30FA movaps    xmm1,xmmword ptr [ebx]
00BE30FD movaps    xmm2,xmmword ptr [esp+40h]
00BE3102 cmpltps   xmm2,xmm1
00BE3106 movmskps  eax,xmm2
00BE3109 cmp      eax,0Ch
00BE310C jne      CompleteBoxPruning+3FEh (0BE315Eh)

00BE310E mov      ecx,dword ptr [esi+4]
00BE3111 cmp      ecx,dword ptr [esi]
00BE3113 jne      CompleteBoxPruning+3BEh (0BE311Eh)
00BE3115 push     1
00BE3117 mov      ecx,esi
00BE3119 call     IceCore::Container::Resize (0BE1240h)
00BE311E mov      edx,dword ptr [esi+4]
00BE3121 mov      eax,dword ptr [esi+8]
00BE3124 mov      ecx,dword ptr [esp+30h]
00BE3128 mov      dword ptr [eax+edx*4],ecx
00BE312B inc      dword ptr [esi+4]
00BE312E mov      eax,dword ptr [esi+4]
00BE3131 mov      edx,dword ptr [edi]
00BE3133 mov      dword ptr [esp+34h],edx
00BE3137 cmp      eax,dword ptr [esi]
00BE3139 jne      CompleteBoxPruning+3E4h (0BE3144h)
00BE313B push     1
00BE313D mov      ecx,esi
00BE313F call     IceCore::Container::Resize (0BE1240h)
00BE3144 mov      eax,dword ptr [esi+4]
00BE3147 mov      ecx,dword ptr [esi+8]
00BE314A mov      edx,dword ptr [esp+34h]
00BE314E movss    xmm0,dword ptr [esp+38h]
00BE3154 mov      dword ptr [ecx+eax*4],edx
00BE3157 inc      dword ptr [esi+4]
00BE315A mov      ecx,dword ptr [esp+3Ch]

00BE315E mov      eax,dword ptr [esp+28h]
00BE3162 add      eax,8
00BE3165 add      ebx,10h
00BE3168 add      edi,4
00BE316B comiss   xmm0,dword ptr [eax]
00BE316E mov      dword ptr [esp+28h],eax
00BE3172 jae      CompleteBoxPruning+39Ah (0BE30FAh)
```

There are still 13 instructions overall in the main `codepath`, with a few twists.

It seems to use three *xmm* registers instead of two, which looks better.

This is the same as in *VC11*, but using two instructions instead of one:

```
00BE30FA movaps    xmm1,xmmword ptr [ebx]
...
00BE3102 cmpltps   xmm2,xmm1
```

xmm2 is still destroyed all the time, and reloaded from the stack:

```
00BE30FD movaps    xmm2,xmmword ptr [esp+40h]
```

The only difference is that it's reloaded in our first block, rather than in the third block before. But it's really all the same otherwise - even the stack offset is the same.

Questions 1 and 2: same.

VC10 seems to manage "*MaxLimit*" better. It's kept in *xmm0* and only reloaded when an overlap occurs:

```
00BE314E movss     xmm0,dword ptr [esp+38h]
```

Well done.

Question 3: *VC10* wins.

In the third block, what used to be *ecx* is now *eax*:

```
00BE3162 add      eax,8
...
00BE316B comiss   xmm0,dword ptr [eax]
00BE316E mov      dword ptr [esp+28h],eax
```

And it's saved to the stack all the time, like in *VC11*.

As for *edi* (previously *edx*), it is not saved to the stack anymore. Yay!

But *eax* (previously *ecx*) is now reloaded from the stack all the time:

```
00BE315E mov      eax,dword ptr [esp+28h]
```

...while it wasn't before. Oh.

Question 4: VC11 wins.

Question 5 and 6: the same overall.

So from just looking at this short snippet it is unclear which compiler does a better job. But then if you look at the code before that inner loop, you find this:

```
0BE2F1B fstp    dword ptr [eax-58h]
00BE2F1E fld     dword ptr [edx+8]
00BE2F21 fstp    dword ptr [eax-54h]
00BE2F24 fld     dword ptr [edx+10h]
00BE2F27 fstp    dword ptr [eax-50h]
00BE2F2A fld     dword ptr [edx+14h]
00BE2F2D fstp    dword ptr [eax-4Ch]
00BE2F30 mov     edx,dword ptr [esi-14h]
00BE2F33 mov     dword ptr [edi-14h],edx
00BE2F36 lea     edx,[edx+edx*2]
00BE2F39 fld     dword ptr [ebx+edx*8]
00BE2F3C lea     edx,[ebx+edx*8]
00BE2F3F fstp    dword ptr [ecx-2Ch]
00BE2F42 fld     dword ptr [edx+0Ch]
00BE2F45 fstp    dword ptr [ecx-28h]
00BE2F48 fld     dword ptr [edx+4]
00BE2F4B fstp    dword ptr [eax-48h]
00BE2F4E fld     dword ptr [edx+8]
00BE2F51 fstp    dword ptr [eax-44h]
00BE2F54 fld     dword ptr [edx+10h]
00BE2F57 fstp    dword ptr [eax-40h]
00BE2F5A fld     dword ptr [edx+14h]
00BE2F5D fstp    dword ptr [eax-3Ch]
00BE2F60 mov     edx,dword ptr [esi-10h]
00BE2F63 mov     dword ptr [edi-10h],edx
00BE2F66 lea     edx,[edx+edx*2]
00BE2F69 fld     dword ptr [ebx+edx*8]
00BE2F6C lea     edx,[ebx+edx*8]
00BE2F6F fstp    dword ptr [ecx-24h]
00BE2F72 fld     dword ptr [edx+0Ch]
00BE2F75 fstp    dword ptr [ecx-20h]
00BE2F78 fld     dword ptr [edx+4]
00BE2F7B fstp    dword ptr [eax-38h]
00BE2F7E fld     dword ptr [edx+8]
00BE2F81 fstp    dword ptr [eax-34h]
00BE2F84 fld     dword ptr [edx+10h]
00BE2F87 fstp    dword ptr [eax-30h]
00BE2F8A fld     dword ptr [edx+14h]
00BE2F8D fstp    dword ptr [eax-2Ch]
```

```

00BE2F90 mov     edx,dword ptr [esi-0Ch]
00BE2F93 mov     dword ptr [edi-0Ch],edx
00BE2F96 lea     edx,[edx+edx*2]
00BE2F99 fld     dword ptr [ebx+edx*8]
00BE2F9C lea     edx,[ebx+edx*8]
00BE2F9F fstp    dword ptr [ecx-1Ch]
00BE2FA2 fld     dword ptr [edx+0Ch]
00BE2FA5 fstp    dword ptr [ecx-18h]
00BE2FA8 fld     dword ptr [edx+4]
00BE2FAB fstp    dword ptr [eax-28h]
00BE2FAE fld     dword ptr [edx+8]
00BE2FB1 fstp    dword ptr [eax-24h]
00BE2FB4 fld     dword ptr [edx+10h]
00BE2FB7 fstp    dword ptr [eax-20h]
00BE2FBA fld     dword ptr [edx+14h]
00BE2FBD fstp    dword ptr [eax-1Ch]

```

Yikes!!

Ok: VC10 still used x87 instructions here and there. VC11 did not.

Forget that inner loop analysis: advantage VC11, big time. The benchmark results confirm this, regardless of the instructions count:

Office PC / unsafe version:

Complete test (brute force): found 11811 intersections in 891760 K-cycles.

19379 K-cycles.

18483 K-cycles.

18943 K-cycles.

18093 K-cycles.

18321 K-cycles.

18417 K-cycles.

18078 K-cycles.

19035 K-cycles.

18500 K-cycles.

18092 K-cycles.

18329 K-cycles.

18092 K-cycles.

17899 K-cycles.

18524 K-cycles.

18085 K-cycles.

19308 K-cycles.

Complete test (box pruning): found 11715 intersections in 17899 K-cycles.

⇒ Much slower than what we got with VC11.

So, VC14 then? Let's have a look:

```
00F32F33 cmpltps   xmm1,xmmword ptr [edi]
00F32F37 movmskps  eax,xmm1
00F32F3A cmp      eax,0Ch
00F32F3D jne      CompleteBoxPruning+2D1h (0F32F91h)

00F32F3F mov      eax,dword ptr [esi+4]
00F32F42 cmp      eax,dword ptr [esi]
00F32F44 jne      CompleteBoxPruning+28Fh (0F32F4Fh)
00F32F46 push     1
00F32F48 mov      ecx,esi
00F32F4A call     IceCore::Container::Resize (0F312D0h)
00F32F4F mov      ecx,dword ptr [esi+4]
00F32F52 mov      eax,dword ptr [esi+8]
00F32F55 mov      edx,dword ptr [esp+34h]
00F32F59 mov      dword ptr [eax+ecx*4],edx
00F32F5C mov      edx,dword ptr [esp+20h]
00F32F60 inc      dword ptr [esi+4]
00F32F63 mov      eax,dword ptr [esi+4]
00F32F66 mov      ecx,dword ptr [edx]
00F32F68 mov      dword ptr [esp+30h],ecx
00F32F6C cmp      eax,dword ptr [esi]
00F32F6E jne      CompleteBoxPruning+2B9h (0F32F79h)
00F32F70 push     1
00F32F72 mov      ecx,esi
00F32F74 call     IceCore::Container::Resize (0F312D0h)
00F32F79 mov      ecx,dword ptr [esi+4]
00F32F7C mov      eax,dword ptr [esi+8]
00F32F7F mov      edx,dword ptr [esp+30h]
00F32F83 mov      dword ptr [eax+ecx*4],edx
00F32F86 inc      dword ptr [esi+4]
00F32F89 mov      ecx,dword ptr [esp+28h]
00F32F8D mov      edx,dword ptr [esp+20h]

00F32F91 movss    xmm0,dword ptr [esp+38h]
00F32F97 add      ecx,8
00F32F9A movaps   xmm1,xmmword ptr [esp+40h]
00F32F9F add      edx,4
00F32FA2 add      edi,10h
00F32FA5 mov      dword ptr [esp+20h],edx
00F32FA9 mov      dword ptr [esp+28h],ecx
00F32FAD comiss   xmm0,dword ptr [ecx]
00F32FB0 jae      CompleteBoxPruning+273h (0F32F33h)
```

This one will be quick: some instructions have been re-ordered in the third block, but other than that it's exactly the same code as *VC11*'s. And pretty much exactly the same performance.

So, nope, a newer compiler doesn't solve the issues.

But are they really issues? There's more to performance than the instructions count, as we saw before. Are we chasing a red herring?

To find our answers, we must go closer to the metal. And that's what we will do next time.

What we learnt:

Different compilers will give different results.

Compilers are getting better all the time, but they're still not perfect.

`std::vector` sucks.